

DASARI.NAGAVENI

EMAIL: dasari.nagaveni2020@vitstudent.ac.in

Vellore institute of technology

Chennai

Assignment-3

Problem Statement: House Price Prediction

Description:- House price prediction is a common problem in the real estate industry and

involves predicting the selling price of a house based on various features and attributes. The

problem is typically approached as a regression problem, where the target variable is the price

of the house, and the features are various attributes of the house

The features used in house price prediction can include both quantitative and categorical

variables, such as the number of bedrooms, house area, bedrooms, furnished, nearness to

main road, and various amenities such as a garage and other factors that may influence the

value of the property.

Accurate predictions can help agents and appraisers price homes correctly, while

homeowners can use the predictions to set a reasonable asking price for their properties.

Accurate house price prediction can also be useful for buyers who are looking to make

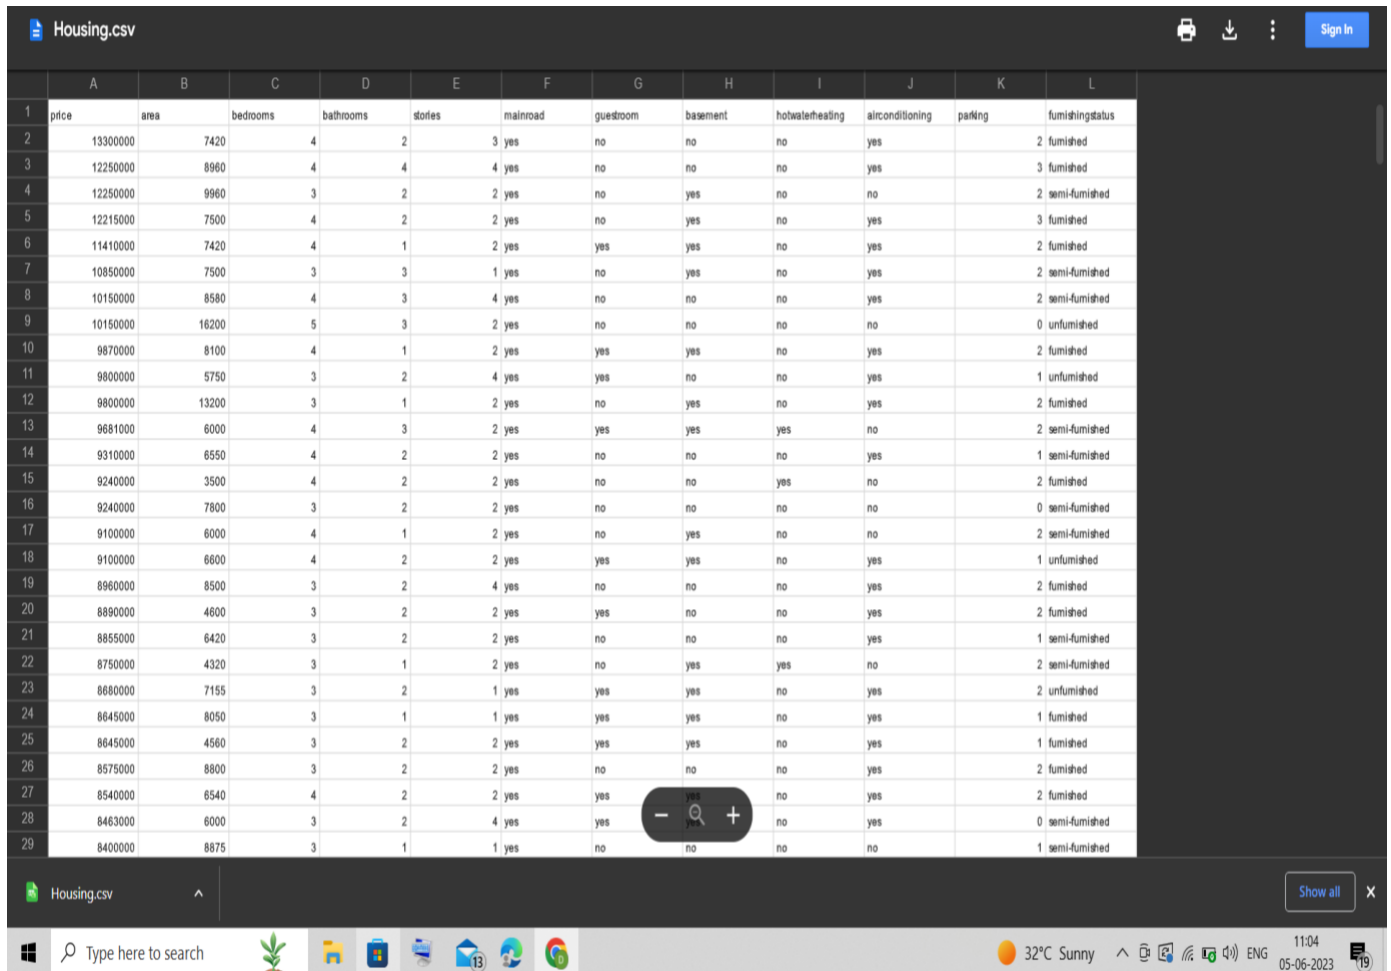
informed decisions about purchasing a property and obtaining a fair price for their investment.

Attribute Information:

Name - Description

- 1- Price-Prices of the houses
- 2- Area- Area of the houses
- 3- Bedrooms- No of house bedrooms
- 4- Bathrooms- No of bathrooms
- 5- Stories- No of house stories
- 6- Main Road- Weather connected to Main road
- 7- Guestroom-Weather has a guest room
- 8- Basement-Weather has a basement
- 9- Hot water heating- Weather has a hot water heater
- 10-Airconditioning-Weather has a air conditioner
- 11-Parking- No of house parking
- 12-Furnishing Status-Furnishing status of house

1. Download the dataset: Dataset



	A	B	C	D	E	F	G	H	I	J	K	L
1	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	furnishingstatus
2	13300000	7420	4	2	3	yes	no	no	no	yes		2 furnished
3	12250000	8960	4	4	4	yes	no	no	no	yes		3 furnished
4	12250000	9960	3	2	2	yes	no	yes	no	no		2 semi-furnished
5	12215000	7500	4	2	2	yes	no	yes	no	yes		3 furnished
6	11410000	7420	4	1	2	yes	yes	yes	no	yes		2 furnished
7	10850000	7500	3	3	1	yes	no	yes	no	yes		2 semi-furnished
8	10150000	8580	4	3	4	yes	no	no	no	yes		2 semi-furnished
9	10150000	16200	5	3	2	yes	no	no	no	no		0 unfurnished
10	9870000	8100	4	1	2	yes	yes	yes	no	yes		2 furnished
11	9800000	5750	3	2	4	yes	yes	no	no	yes		1 unfurnished
12	9800000	13200	3	1	2	yes	no	yes	no	yes		2 furnished
13	9681000	6000	4	3	2	yes	yes	yes	yes	no		2 semi-furnished
14	9310000	6550	4	2	2	yes	no	no	no	yes		1 semi-furnished
15	9240000	3500	4	2	2	yes	no	no	yes	no		2 furnished
16	9240000	7800	3	2	2	yes	no	no	no	no		0 semi-furnished
17	9100000	6000	4	1	2	yes	no	yes	no	no		2 semi-furnished
18	9100000	6600	4	2	2	yes	yes	yes	no	yes		1 unfurnished
19	8960000	8500	3	2	4	yes	no	no	no	yes		2 furnished
20	8890000	4600	3	2	2	yes	yes	no	no	yes		2 furnished
21	8855000	6420	3	2	2	yes	no	no	no	yes		1 semi-furnished
22	8750000	4320	3	1	2	yes	no	yes	yes	no		2 semi-furnished
23	8680000	7155	3	2	1	yes	yes	yes	no	yes		2 unfurnished
24	8645000	8050	3	1	1	yes	yes	yes	no	yes		1 furnished
25	8645000	4560	3	2	2	yes	yes	yes	no	yes		1 furnished
26	8575000	8800	3	2	2	yes	no	no	no	yes		2 furnished
27	8540000	8540	4	2	2	yes	yes	yes	no	yes		2 furnished
28	8463000	6000	3	2	4	yes	yes	yes	no	yes		0 semi-furnished
29	8400000	8875	3	1	1	yes	no	no	no	no		1 semi-furnished

2. Load the dataset into the tool.

```
✓ [6] import numpy as np
1s      import pandas as pd
      import difflib
      from sklearn.feature_extraction.text import TfidfVectorizer
      from sklearn.metrics.pairwise import cosine_similarity
```

```
✓ [7] df=pd.read_csv('/content/Housing.csv')
1s
```

```
import numpy as np
import pandas as pd
import difflib
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
df=pd.read_csv('/content/Housing.csv')
```

3. Perform Below Visualizations.

🔍 Univariate Analysis

🔍 Bi-Variate Analysis

🔍 Multi-Variate Analysis

Statistics summary

```
summary_stats = df.describe()
print(summary_stats)
```

```
✓ 1s [7] df=pd.read_csv('/content/Housing.csv')

✓ 0s summary_stats = df.describe()
print(summary_stats)
```

	price	area	bedrooms	bathrooms	stories	\
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	

	parking
count	545.000000
mean	0.693578
std	0.861586
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

Univariate Analysis

HISTOGRAM

Histogram

```
import matplotlib.pyplot as plt
```


```
plt.hist(df['price'], bins=20)
```

```
plt.xlabel('Price')
```

```
plt.ylabel('Frequency')
```

```
plt.title('Distribution of Price')
```

```
plt.show()
```



BOXPLOT

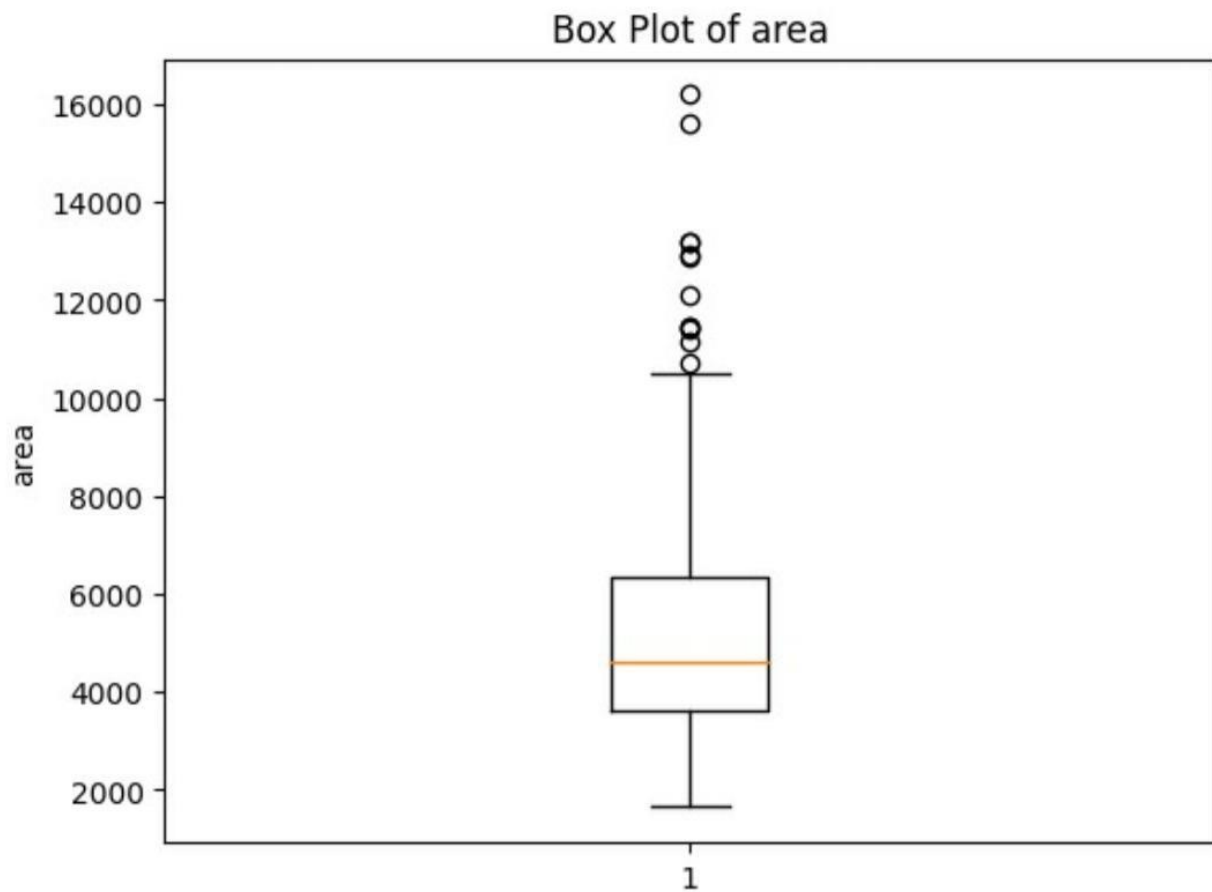
Box plot

```
plt.boxplot(df['area'])
```

```
plt.ylabel('area')
```

```
plt.title('Box Plot of area')
```

```
plt.show()
```



COUNT OF UNIQUE VALUES

Count of unique values

```
value_counts = df['bedroom
```

```
(kind='bar')
```

```
plt.xlabel('Bedrooms')
```

```
plt.ylabel('Count')
```

```
plt.title('Number of Bedrooms')
```

```
plt.show()
```



```

3  300
2  136
4   95
5   10
6    2
1    2
2  Name: bedrooms, dtype: int64

```

Bar chart

Bar chart

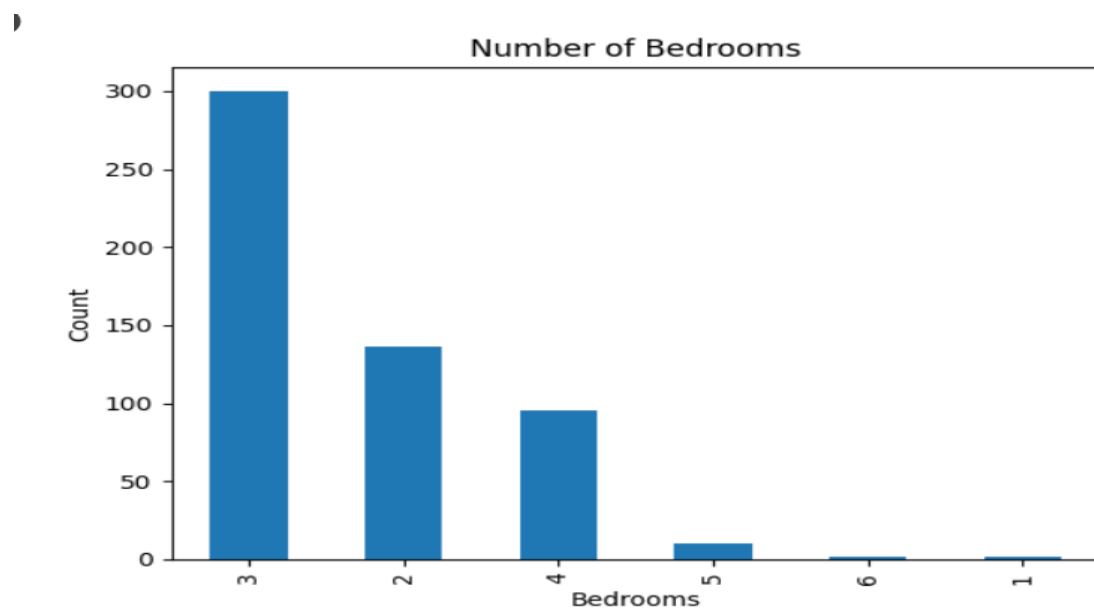
```
Value_counts.plot(kind='bar')
```

```
Plt.xlabel('Bedrooms')
```

```
Plt.ylabel('Count')
```

```
Plt.title('Number of Bedrooms')
```

```
Plt.show()
```




DENSITY PLOT

Density plot

```
df['price'].plot(kind='density')
```

```
plt.xlabel('Price')
plt.title('Density Plot of Price')
plt.show()
```




BI-VARIATE ANALYSIS

SCATTER PLOT

```
# Scatter plot
import seaborn as sns

sns.scatterplot(x='area', y='price', data=df)

plt.xlabel('area')
plt.ylabel('Price')
plt.title('Scatter Plot: area vs Price')
plt.show()
```



BOX PLOT

```
# Box plot
```

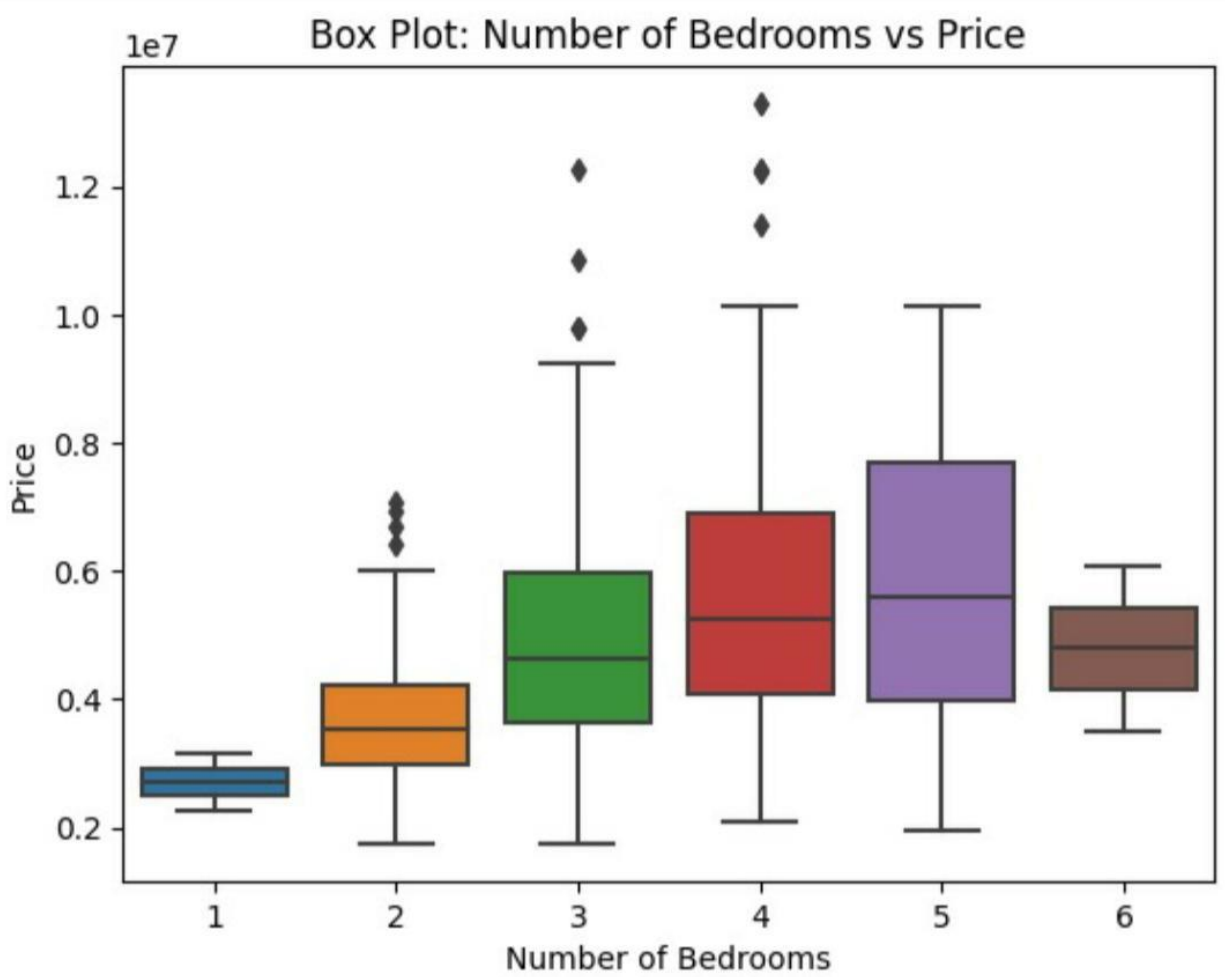
```
sns.boxplot(x='bedrooms', y='price', data=df)
```

```
plt.xlabel('Number of Bedrooms')
```

```
plt.ylabel('Price')
```

```
plt.title('Box Plot: Number of Bedrooms vs Price')
```

```
plt.show()
```



HEAT MAP


```
# Heatmap
```

```
correlation_matrix = df.corr()
```

```
sns.heatmap(correlation_matrix, annot=True)
```

```
plt.title('Correlation Matrix Heatmap')
```

```
plt.show()
```



PAIRPLOT

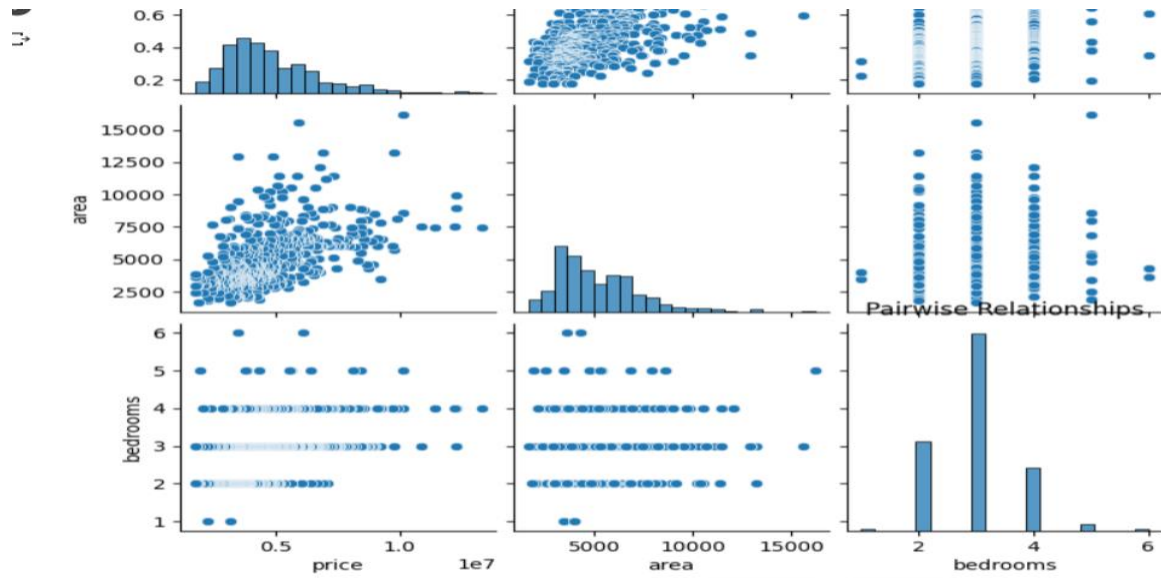
```
variables_of_interest = ['price', 'area', 'bedrooms']
```

```
# Pairplot
```

```
sns.pairplot(df[variables_of_interest])
```

```
plt.title('Pairwise Relationships')
```

```
plt.show()
```



CORRELATION MATRIX


Correlation matrix heatmap

```
corr_matrix = df[variables_of_interest].corr()
```

```
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
```

```
plt.title('Correlation Matrix Heatmap')
```

```
plt.show()
```



4. Perform descriptive statistics on the dataset.

```
# Calculate descriptive statistics
```

```
descriptive_stats = df.describe()
```

```
# Print the descriptive statistics
```

```
print(descriptive_stats)
```

	price	area	bedrooms	bathrooms	stories	\
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	
	parking					
count	545.000000					
mean	0.693578					
std	0.861586					
min	0.000000					
25%	0.000000					
50%	0.000000					
75%	1.000000					
max	3.000000					

5. Check for Missing values and deal with them.

```
# Check for missing values
```

```
missing_values = df.isnull().sum()
```

```
print(missing_values)
```

```
# Option 1: Drop rows with missing values
```

```
df_dropped = df.dropna()
```

```
print("Shape after dropping missing values:", df_dropped.shape)
```


Option 2: Fill missing values with mean/median/mode

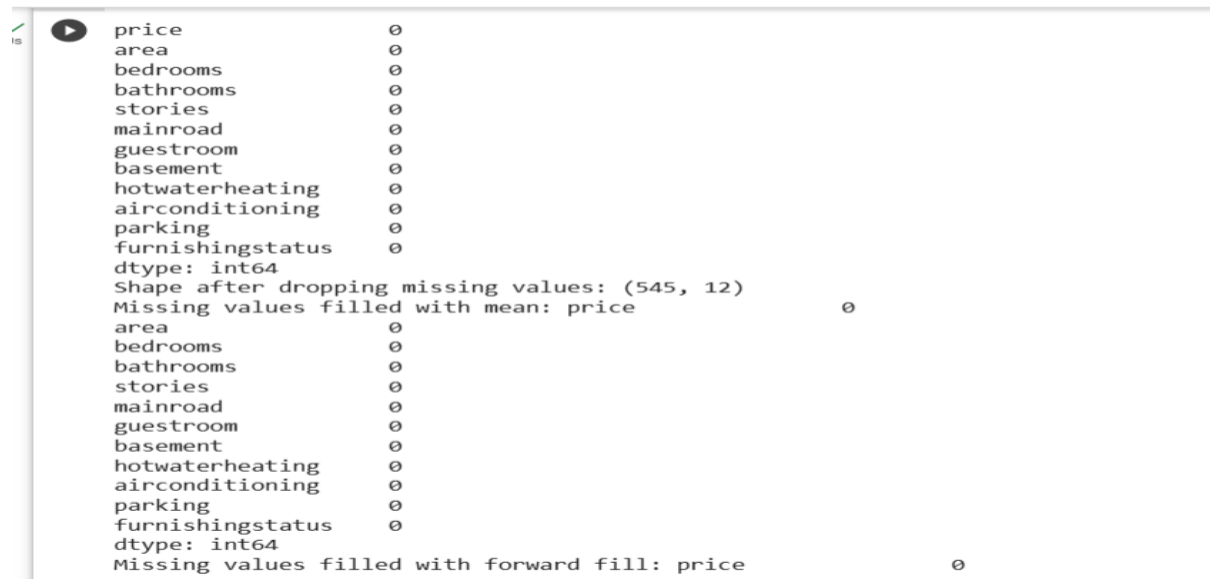
For example, fill missing values in the 'price' column with the mean

```
mean_price = df['price'].mean()
df_filled = df.fillna({'price': mean_price})
print("Missing values filled with mean:", df_filled.isnull().sum())
```

Option 3: Fill missing values with forward fill or backward fill

For example, fill missing values using forward fill method

```
df_ffilled = df.fillna(method='ffill')
print("Missing values filled with forward fill:", df_ffilled.isnull().sum())
```



```
price      0
area       0
bedrooms   0
bathrooms  0
stories    0
mainroad   0
guestroom  0
basement   0
hotwaterheating  0
airconditioning  0
parking    0
furnishingstatus  0
dtype: int64
Shape after dropping missing values: (545, 12)
Missing values filled with mean: price      0
area      0
bedrooms  0
bathrooms 0
stories   0
mainroad  0
guestroom 0
basement  0
hotwaterheating 0
airconditioning 0
parking    0
furnishingstatus 0
dtype: int64
Missing values filled with forward fill: price      0
-----
```

```
+ Code + Text
0s ▶ missing values filled with mean: price
area 0
bedrooms 0
bathrooms 0
stories 0
mainroad 0
guestroom 0
basement 0
hotwaterheating 0
airconditioning 0
parking 0
furnishingstatus 0
dtype: int64
Missing values filled with forward fill: price 0
area 0
bedrooms 0
bathrooms 0
stories 0
mainroad 0
guestroom 0
basement 0
hotwaterheating 0
airconditioning 0
parking 0
furnishingstatus 0
dtype: int64
```

6. Find the outliers and replace them outliers

Define the columns to check for outliers

```
columns_to_check = ['price', 'area']
```

Define the outlier detection method (e.g., z-score or IQR)

```
outlier_method = 'z-score'
```

Define the threshold for outlier detection

```
threshold = 3
```

Function to replace outliers with a suitable value (e.g., median)

```
def replace_outliers(data, col):
    if outlier_method == 'z-score':
        z_scores = (data[col] - data[col].mean()) / data[col].std()
        outliers = np.abs(z_scores) > threshold

    elif outlier_method == 'iqr':
        Q1 = data[col].quantile(0.25)
        Q3 = data[col].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
outliers = (data[col] < (Q1 - 1.5 * IQR)) | (data[col] > (Q3 + 1.5 * IQR))
```

```
else:
```

```
    print("Invalid outlier detection method")
```

```
    return data
```

```
data.loc[outliers, col] = data[col].median()
```

```
return data
```

```
# Iterate over the columns to check for outliers and replace them
```

```
for col in columns_to_check:
```

```
    df = replace_outliers(df, col)
```

```
# Display the modified dataset
```

```
print(df)
```

```
0  4340000 7420 4 2 3 yes no no
1  4340000 8960 4 4 yes no no
2  4340000 9960 3 2 yes no yes
3  4340000 7500 4 2 yes no yes
4  4340000 7420 4 1 2 yes yes yes
..  ...  ...  ...  ...  ...  ...  ...  ...
540 1820000 3000 2 1 1 yes no yes
541 1767150 2400 3 1 1 no no no
542 1750000 3620 2 1 1 yes no no
543 1750000 2910 3 1 1 no no no
544 1750000 3850 3 1 2 yes no no

hotwaterheating airconditioning parking furnishingstatus
0 no yes 2 furnished
1 no yes 3 furnished
2 no no 2 semi-furnished
3 no yes 3 furnished
4 no yes 2 furnished
..  ...  ...  ...
540 no no 2 unfurnished
541 no no 0 semi-furnished
542 no no 0 unfurnished
543 no no 0 furnished
544 no no 0 unfurnished

[545 rows x 12 columns]
```

7. Check for Categorical columns and perform encoding.

```
# Identify categorical columns
```

```
categorical_columns = df.select_dtypes(include=['object']).columns
```

```
print("Categorical Columns:", categorical_columns)
```

```
# Perform one-hot encoding
```

```
df_encoded = pd.get_dummies(df, columns=categorical_columns)
```

```
print("Encoded DataFrame:")
```

```
print(df_encoded)
```

```
Categorical Columns: Index(['mainroad', 'guestroom', 'basement', 'hotwaterheating',  
                             'airconditioning', 'furnishingstatus'],  
                             dtype='object')
```

```
Encoded DataFrame:
```

	price	area	bedrooms	bathrooms	stories	parking	mainroad_no	\
0	4340000	7420	4	2	3	2	0	
1	4340000	8960	4	4	4	3	0	
2	4340000	9960	3	2	2	2	0	
3	4340000	7500	4	2	2	3	0	
4	4340000	7420	4	1	2	2	0	
..	
540	1820000	3000	2	1	1	2	0	
541	1767150	2400	3	1	1	0	1	
542	1750000	3620	2	1	1	0	0	
543	1750000	2910	3	1	1	0	1	
544	1750000	3850	3	1	2	0	0	

	mainroad_yes	guestroom_no	guestroom_yes	basement_no	basement_yes	\
0	1	1	0	1	0	
1	1	1	0	1	0	
2	1	1	0	0	1	
3	1	1	0	0	1	
4	1	0	1	0	1	
..	
540	1	1	0	0	1	
541	0	1	0	1	0	
542	1	1	0	1	0	
543	0	1	0	1	0	
544	1	1	0	1	0	

	hotwaterheating_no	hotwaterheating_yes	airconditioning_no	\
0	1	0	0	

1	1	0	0
2	1	0	1
3	1	0	0
4	1	0	0
..
540	1	0	1
541	1	0	1
542	1	0	1
543	1	0	1
544	1	0	1

	airconditioning_yes	furnishingstatus_furnished \
0	1	1
1	1	1
2	0	0
3	1	1
4	1	1
..
540	0	0
541	0	0
542	0	0
543	0	1
544	0	0

	furnishingstatus_semi-furnished	furnishingstatus_unfurnished
0	0	0
1	0	0
2	1	0
3	0	0
4	0	0
..
540	0	1
541	1	0
542	0	1
543	0	0
544	0	1

[545 rows x 19 columns]

8. Split the data

into dependent and independent variables.

Split the data into dependent and independent variables

X = df.drop('price', axis=1) # Independent variables

y = df['price'] # Dependent variable

```
# Print the shape of the variables
```

```
print("Shape of X:", X.shape)
```

```
print("Shape of y:", y.shape)
```

```
|  ─  Shape of y: (545,)
```

9. Scale the independent variables

```
import pandas as pd
```

```
from sklearn.preprocessing import StandardScaler
```

```
#Assuming your Titanic dataset is stored in a pandas DataFrame called 'df'
```

```
#Assuming your independent variables are stored in a DataFrame called
```

```
#Select the independent variables you want to scale
```

```
independent_variables = ['price', 'area', 'bathrooms']
```

```
# Create a new DataFrame with only the selected independent variables
```

```
X_selected = df[independent_variables]
```

```
#Initialize the StandardScaler
```

```
scaler = StandardScaler()
```

```
# Fit and transform the selected independent variables
```

```
X_scaled = scaler.fit_transform(X_selected)
```

```
# Create a new DataFrame with the scaled independent variables
```

```
X_scaled_df = pd.DataFrame(X_scaled, columns=independent_variables)
```

```
# Print the first few rows of the scaled independent variables
```

```
print(X_scaled_df.head())
```

```

import pandas as pd

from sklearn.model_selection import train_test_split


# Split the data into dependent and independent variables
X = df.drop('price', axis=1) # Independent variables
y = df['price']             # Dependent variable


# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Print the shapes of the train and test sets
print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)

```

```

Shape of y_test: (109,)

```

11. Build the Model

```

import pandas as pd

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

```



```
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.metrics import mean_squared_error
```

```
# Split the data into dependent and independent variables
```

```
X = df.drop('price', axis=1) # Independent variables
```

```
y = df['price'] # Dependent variable
```

```
# Identify the categorical columns
```

```
categorical_cols = X.select_dtypes(include=['object']).columns
```

```
# Perform one-hot encoding for categorical variables
```

```
ct = ColumnTransformer([('encoder', OneHotEncoder(), categorical_cols)], remainder='passthrough')
```

```
X_encoded = ct.fit_transform(X)
```

```
# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)
```

```
# Build the linear regression model
```

```
lm = LinearRegression()
```

```
# Fit the model to the training data
```

```
lm.fit(X_train, y_train)
```

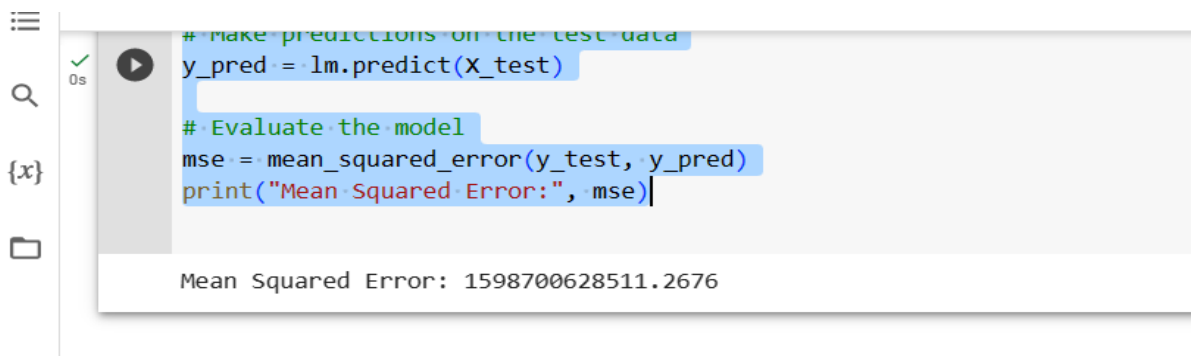
```
# Make predictions on the test data
```

```
y_pred = lm.predict(X_test)
```

```
# Evaluate the model
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
print("Mean Squared Error:", mse)
```

A screenshot of a Jupyter Notebook interface. On the left is a sidebar with icons for a menu, search, variables, and files. The main area shows a code cell with the following Python code:

```
# Make predictions on the test data
y_pred = lm.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

 The code is highlighted with blue and green background colors. Below the code cell, the output is displayed:

```
Mean Squared Error: 1598700628511.2676
```

12. Train the Model

```
# Fit the model to the training data
```

```
lm.fit(X_train, y_train)
```

13. Test the Model

```
# Make predictions on the test data
```

```
y_pred = model.predict(X_test)
```

```
y_pred
```

```
+ Code + Text
array([4911760.61705211, 6434232.18391202, 3392953.92431767,
4840490.1994299 , 3572716.04843916, 3911933.17546033,
5627409.12656888, 5830247.47683143, 2762955.4117319 ,
2681995.10143281, 8431281.03976994, 3067236.35258577,
3121790.14209741, 3469409.9747259 , 3962974.95809195,
4833083.00985176, 3259802.94309297, 4905726.89310817,
4548623.99581589, 3885979.44593909, 5306463.18766681,
5613667.34590536, 2995546.41704388, 4364278.85000292,
5296878.64523837, 6774708.43116889, 3549041.57237742,
5185779.38620721, 5384144.85802853, 3567335.93137482,
5732174.4579619 , 3499121.13278727, 6639779.66999354,
4564573.07160303, 3946575.09941465, 5633058.90120369,
4988885.71729119, 4590311.84353848, 3205315.93006053,
4581988.80582526, 4753426.11379039, 3672927.18390534,
6414108.6932336 , 4125831.55833053, 4074637.17104933,
4352905.0841844 , 6632469.04344548, 4324417.66225873,
4102589.94799076, 3420577.35610439, 6929935.96980785,
2945041.6056961 , 4499948.03878172, 4503731.15749593,
4023861.81904184, 2734819.26446367, 6714595.49809504,
3205262.7099005 , 4666964.63207716, 3033529.50569941,
4545348.3050988 , 3345266.82659203, 5017059.76091084,
4370268.96006151, 4175325.87856412, 4871563.03799354,
6376116.96606657, 3718977.63361604, 5956657.85327391,
5612164.12660341, 3887909.26815781, 4880427.47614069,
4658420.52787432, 7155946.40609129, 3547713.9134139 ,
5696176.24955518, 4039209.33922616, 4270060.3361625 ,
4880772.16341754, 3663125.34299471, 6768406.34878505,
4152403.9742923 , 5660134.45239474, 5022140.79956395,
3029182.16646428, 6468491.39509035, 2925882.13221748,
3881184.47841876, 7035042.60027418, 7833903.38020682,
3456607.8527137 , 6090544.65850122, 3803619.80232645,
3866759.85337504, 7189010.49652199, 5044688.39352709,
5601819.85702438, 6283106.37016643, 4477793.01154557,
5467926.93803639, 3941107.32216884, 6133114.5650784,
4027694.55864242, 5654372.96701307, 5142052.75420025,
4382665.96890903, 6889112.31022481, 6316539.29089781,
```

```
array([4911760.61705211, 6434232.18391202, 3392953.92431767,
4840490.1994299 , 3572716.04843916, 3911933.17546033,
5627409.12656888, 5830247.47683143, 2762955.4117319 ,
2681995.10143281, 8431281.03976994, 3067236.35258577,
3121790.14209741, 3469409.9747259 , 3962974.95809195,
4833083.00985176, 3259802.94309297, 4905726.89310817,
4548623.99581589, 3885979.44593909, 5306463.18766681,
5613667.34590536, 2995546.41704388, 4364278.85000292,
5296878.64523837, 6774708.43116889, 3549041.57237742,
5185779.38620721, 5384144.85802853, 3567335.93137482,
5732174.4579619 , 3499121.13278727, 6639779.66999354,
4564573.07160303, 3946575.09941465, 5633058.90120369,
4988885.71729119, 4590311.84353848, 3205315.93006053,
4581988.80582526, 4753426.11379039, 3672927.18390534,
6414108.6932336 , 4125831.55833053, 4074637.17104933,
4352905.0841844 , 6632469.04344548, 4324417.66225873,
4102589.94799076, 3420577.35610439, 6929935.96980785,
2945041.6056961 , 4499948.03878172, 4503731.15749593,
4023861.81904184, 2734819.26446367, 6714595.49809504,
3205262.7099005 , 4666964.63207716, 3033529.50569941,
4545348.3050988 , 3345266.82659203, 5017059.76091084,
4370268.96006151, 4175325.87856412, 4871563.03799354,
6376116.96606657, 3718977.63361604, 5956657.85327391,
5612164.12660341, 3887909.26815781, 4880427.47614069,
4658420.52787432, 7155946.40609129, 3547713.9134139 ,
5696176.24955518, 4039209.33922616, 4270060.3361625 ,
4880772.16341754, 3663125.34299471, 6768406.34878505,
4152403.9742923 , 5660134.45239474, 5022140.79956395,
3029182.16646428, 6468491.39509035, 2925882.13221748,
3881184.47841876, 7035042.60027418, 7833903.38020682,
```

```
3456607.8527137 , 6090544.65850122, 3803619.80232645,  
3866759.85337504, 7189010.49652199, 5044688.39352709,  
5601819.85702438, 6283106.37016643, 4477793.01154557,  
5467926.03803639, 3941107.32216884, 6133114.56560784,  
4027694.55864242, 5654372.96701307, 5142052.75420025,  
4382665.96890903, 6889112.31022481, 6316539.29089781,  
5479741.32966125])
```

14. Measure the performance using Metrics.

```
from sklearn.metrics import r2_score
```

```
r2 = r2_score(y_test, y_pred)
```

```
print("R-squared Score:", r2)
```

```
r2
```

```
R-squared Score: 0.5928639479880735
```

```
0.5928639479880735
```